

Python 3.5.6 |Anaconda 4.2.0 (64-bit)| (default, Aug 26 2018, 16:05:27) [MSC v.1900 64 bit (AMD64)]
Type "copyright", "credits" or "license" for more information.

IPython 5.1.0 -- An enhanced Interactive Python.

? -> Introduction and overview of IPython's features.

%quickref -> Quick reference.

help -> Python's own help system.

object? -> Details about 'object', use 'object??' for extra details.

In [1]:

```
In [1]: import pandas as pd
...: import matplotlib.pyplot as plt
...: import numpy as np
```

```
In [2]: phishingData = pd.read_csv('phishing.txt')
...: X = phishingData.iloc[:, :-1].values
...: y = phishingData.iloc[:, 30].values
```

```
In [3]: from sklearn.cross_validation import train_test_split
...: X_train, X_test, y_train, y_test = train_test_split
(X, y, test_size=0.3, random_state=4)
...:
```

```
In [4]: from sklearn.preprocessing import StandardScaler
...: scalar = StandardScaler()
...: X_train = scalar.fit_transform(X_train)
...: X_test = scalar.fit_transform(X_test)
...:
```

C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.

warnings.warn(msg, DataConversionWarning)

C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.

warnings.warn(msg, DataConversionWarning)

C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.

warnings.warn(msg, DataConversionWarning)

C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.

warnings.warn(msg, DataConversionWarning)

```
In [5]: from sklearn.linear_model import LogisticRegression
...: LRclassifier = LogisticRegression(C=100, random_state=0)
...: LRclassifier.fit(X_train, y_train)
...:
...: LRpredict = LRclassifier.predict(X_test)
```

```
In [6]: LRclassifier.score(X_train, y_train)
...:
```

Out[6]: 0.92981775882124851

```
In [7]: LRclassifier.score(X_test, y_test)
```

Out[7]: 0.92734398552909258

```
In [8]: from sklearn.metrics import confusion_matrix
...: confusionMatrix = confusion_matrix(y_test, LRpredict)
```

```

....:

In [9]: confusionMatrix
Out[9]:
array([[1357, 153],
       [ 88, 1719]])

In [10]: X = phishingData.iloc[0:5,[6,14]].values
....: y = phishingData.iloc[0:5,30].values
....:

In [11]: from sklearn.cross_validation import train_test_split
....: X_train,X_test,y_train,y_test = train_test_split
(X,y,test_size=0.3,random_state=4)
....:

In [12]: from sklearn.preprocessing import StandardScaler
....: scalar = StandardScaler()
....: X_train = scalar.fit_transform (X_train)
....: X_test = scalar.fit_transform (X_test)
....:
C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.
  warnings.warn(msg, DataConversionWarning)
C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.
  warnings.warn(msg, DataConversionWarning)
C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.
  warnings.warn(msg, DataConversionWarning)
C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.
  warnings.warn(msg, DataConversionWarning)

In [13]: from sklearn.linear_model import LogisticRegression
....: LRclassifier1 = LogisticRegression(C=100,random_state=0)
....: LRclassifier1.fit(X_train,y_train)
....:
Out[13]:
LogisticRegression(C=100, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=0, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False)

In [14]: LRpredict1 = LRclassifier1.predict(X_test)
....:

In [15]: LRclassifier1.score(X_train,y_train)
....:
Out[15]: 1.0

In [16]: LRclassifier1.score(X_test,y_test)
....:
Out[16]: 1.0

In [17]: from sklearn.metrics import confusion_matrix
....: LRconfusionMatrix1 = confusion_matrix(y_test,LRpredict1)

```

```
In [18]: LRconfusionMatrix1
```

```
Out[18]:
```

```
array([[1, 0],  
       [0, 1]])
```

```
In [19]: xx, yy = np.mgrid[-5:5:.01, -5:5:.01]
```

```
....: grid = np.c_[xx.ravel(), yy.ravel()]
```

```
....: probs = LRclassifier1.predict_proba(grid)[: , 1].reshape(xx.shape)
```

```
....:
```

```
....: print(probs)
```

```
....:
```

```
....: f, ax = plt.subplots(figsize=(8, 6))
```

```
....: contour = ax.contourf(xx, yy, probs, 25, cmap="RdBu",
```

```
....:                      vmin=0, vmax=1)
```

```
....: ax_c = f.colorbar(contour)
```

```
....: ax_c.set_label("$P(y = 1)$")
```

```
....: ax_c.set_ticks([0, .25, .5, .75, 1])
```

```
....:
```

```
....: ax.scatter(X_test[:, 0], X_test[:, 1], c = (y_test == 1 ), s=50,
```

```
....:           cmap="RdBu", vmin=-.2, vmax=1.2,
```

```
....:           edgecolor="white", linewidth=1)
```

```
....:
```

```
....: ax.set(aspect="equal",
```

```
....:       xlim=(-5, 5), ylim=(-5, 5),
```

```
....:       xlabel="$X_1$", ylabel="$X_2$")
```

```
....:
```

```
....: plt.show()
```

```
....:
```

```
[[ 8.12989106e-11  8.17055455e-11  8.21142142e-11 ...,  1.17595783e-08
```

```
 1.18183965e-08  1.18775089e-08]
```

```
[ 8.44905622e-11  8.49131608e-11  8.53378731e-11 ...,  1.22212386e-08
```

```
 1.22823659e-08  1.23437989e-08]
```

```
[ 8.78075123e-11  8.82467013e-11  8.86880871e-11 ...,  1.27010228e-08
```

```
 1.27645498e-08  1.28283946e-08]
```

```
...,
```

```
[ 9.99999739e-01  9.99999740e-01  9.99999742e-01 ...,  9.99999998e-01
```

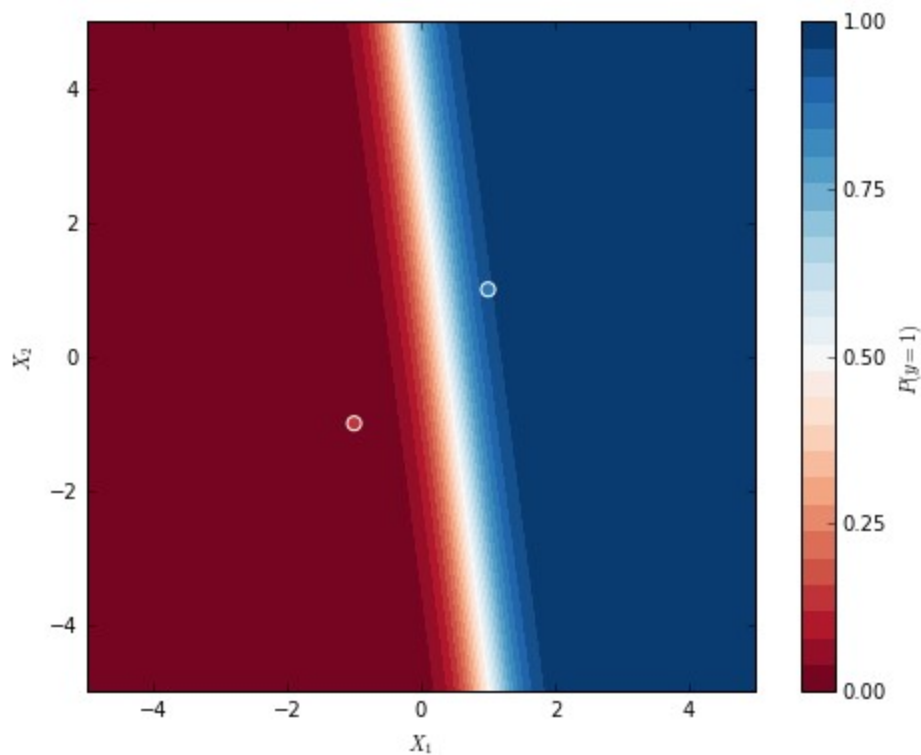
```
 9.99999998e-01  9.99999998e-01]
```

```
[ 9.99999749e-01  9.99999750e-01  9.99999751e-01 ...,  9.99999998e-01
```

```
 9.99999998e-01  9.99999998e-01]
```

```
[ 9.99999758e-01  9.99999760e-01  9.99999761e-01 ...,  9.99999998e-01
```

```
 9.99999998e-01  9.99999998e-01]]
```



```
In [20]: X = phishingData.iloc[0:13,[6,14]].values
...: y = phishingData.iloc[0:13,30].values
...:
```

```
In [21]: from sklearn.cross_validation import train_test_split
...: X_train,X_test,y_train,y_test = train_test_split
(X,y,test_size=0.3,random_state=4)
...:
...: #perform feature scaling
```

```
In [22]: from sklearn.preprocessing import StandardScaler
...: scalar = StandardScaler()
...: X_train = scalar.fit_transform (X_train)
...: X_test = scalar.transform (X_test)
C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.
warnings.warn(msg, DataConversionWarning)
C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.
warnings.warn(msg, DataConversionWarning)
C:\Users\SRIKANT\Anaconda3\lib\site-packages\sklearn\utils\validation.py:420:
DataConversionWarning: Data with input dtype int64 was converted to float64 by
StandardScaler.
warnings.warn(msg, DataConversionWarning)
```

```
In [23]: from sklearn.linear_model import LogisticRegression
...: LRclassifier11 = LogisticRegression(C=100,random_state=0)
...: LRclassifier11.fit(X_train,y_train)
```

```
Out[23]:
LogisticRegression(C=100, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=0, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False)
```

```

In [24]: LRpredict11 = LRclassifier11.predict(X_test)

In [25]: LRclassifier11.score(X_train,y_train)
....:
Out[25]: 0.88888888888888884


In [26]: LRclassifier11.score(X_test,y_test)
....:
Out[26]: 1.0

In [27]: from sklearn.metrics import confusion_matrix
....: LRconfusionMatrix11 = confusion_matrix(y_test,LRpredict11)
....:

In [28]: LRconfusionMatrix11
Out[28]:
array([[2, 0],
       [0, 2]])

In [29]: xx, yy = np.mgrid[-5:5:.01, -5:5:.01]
....: grid = np.c_[xx.ravel(), yy.ravel()]
....: probs = LRclassifier11.predict_proba(grid)[: , 1].reshape(xx.shape)
....:
....: print(probs)
....:
....: f, ax = plt.subplots(figsize=(8, 6))
....: contour = ax.contourf(xx, yy, probs, 25, cmap="RdBu",
....:                      vmin=0, vmax=1)
....: ax_c = f.colorbar(contour)
....: ax_c.set_label("$P(y = 1)$")
....: ax_c.set_ticks([0, .25, .5, .75, 1])
....:
....: ax.scatter(X_test[:, 0], X_test[:, 1],c = (y_test == 1 ), s=50,
....:          cmap="RdBu", vmin=-.2, vmax=1.2,
....:          edgecolor="white", linewidth=1)
....:
....: ax.set(aspect="equal",
....:       xlim=(-5, 5), ylim=(-5, 5),
....:       xlabel="$X_1$", ylabel="$X_2$")
....:
....: plt.show()
....:
[[ 5.10217931e-07  5.15342853e-07  5.20519253e-07 ...,  1.07298055e-02
  1.08364140e-02  1.09440700e-02]
 [ 5.18016238e-07  5.23219491e-07  5.28475009e-07 ...,  1.08920165e-02
  1.10002188e-02  1.11094838e-02]
 [ 5.25933737e-07  5.31216518e-07  5.36552362e-07 ...,  1.10566524e-02
  1.11664717e-02  1.12773693e-02]
 ...,
 [ 6.53557676e-01  6.55817138e-01  6.58069573e-01 ...,  9.99975065e-01
  9.99975313e-01  9.99975558e-01]
 [ 6.56984116e-01  6.59232882e-01  6.61474502e-01 ...,  9.99975440e-01
  9.99975684e-01  9.99975926e-01]
 [ 6.60394275e-01  6.62632166e-01  6.64862792e-01 ...,  9.99975810e-01
  9.99976050e-01  9.99976288e-01]]

```



In [30]: